



















When to use what?

- **Optimal Transport**

- ✓ When geometry of support matters (images, shapes, distributions on a metric space).
- ✓ When supports may not overlap (OT still finite).
- ✗ Heavy computational/statistical cost in high dimension (except in some cases + mitigated by EOT)

- **f-Divergences** (KL, JS, χ^2 , ...)

- ✓ When densities overlap and likelihood ratios are meaningful (language models, inference).
- ✓ Strong information-theoretic foundations (variational inference, GANs).
- ✗ KL breaks down if supports are disjoint (divergence = ∞).

- **Integral Probability Metrics** (IPMs: MMD, Energy, ...)

- ✓ When efficient, dimension-free estimation is needed (two-sample tests, kernels).
- ✓ Flexible via choice of function class $\mathcal{F}$.
- ✗ Geometry depends heavily on $\mathcal{F}$; may miss fine structure.

Takeaways & Preview

- OT is part of a broader family of statistical divergences, two close cousins:
 - f -divergences (likelihood ratio-based)
 - Integral Probability Metrics (witness-based)
- Choosing the right divergence = trade-off between geometry, smoothness, and tractability
- **Next time:** deeper dive into OT as an ML loss function (differentiability, mini-batching, backprop, etc)