

Application: Knowledge Distillation

Sinkhorn Distance Minimization for Knowledge Distillation

Xiao Cui¹, Yulei Qin^{2*}, Yuting Gao², Enwei Zhang², Zihan Xu², Tong Wu²,
Ke Li², Xing Sun², Wengang Zhou^{1†} and Houqiang Li^{1†}


¹University of Science and Technology of China, ²Tencent YouTu Lab

cuixiao@mail.ustc.edu.cn, {zhwg,lihq}@ustc.edu.cn,

{yuleiqin, yutinggao, miyozhang, ianxxu, townswu, tristanli, winfredsun}@tencent.com

Abstract

Knowledge distillation (KD) has been widely adopted to compress large language models (LLMs). Existing KD methods investigate various divergence measures including the Kullback-Leibler (KL), reverse Kullback-Leibler



3.1. Problem Statement

Given a sample x_i and its ground-truth label $y_i \in \mathbb{R}^d$ in the training set, the output logits with softmax activation σ_τ from the teacher f_T and the student f_S are respectively $t_i \in \mathbb{R}^d$ and $s_i \in \mathbb{R}^d$:

$$t_i = \sigma_\tau(f_T(x_i)), \quad s_i = \sigma_\tau(f_S(x_i)), \quad (1)$$

where τ is the temperature and d is the dimension of the output logits. The objective of KD is to minimize the measured divergence $J(t_i, s_i)$.

Method	#Params.	COLA (MCC)	SST-2 (ACC)	MNLI-(m/mm) (ACC)	MRPC (F1)	RTE (ACC)	QNLI (ACC)	QQP (ACC)	Avg
BERT _{base} (T) (Devlin et al., 2018)	110M	60.3	93.7	84.9/84.8	91.4	71.1	91.7	91.5	83.28
BERT ₆ (S) (Turc et al., 2019)	66M	51.2	91.0	81.7/82.6	89.2	66.1	89.3	90.4	79.53
Task-specific Representation-based Distillation									
PKD (Sun et al., 2019)	66M	45.5	91.3	81.3/-	85.7	66.5	88.4	88.4	77.63
TinyBERT (Jiao et al., 2020)	66M	53.8	92.3	83.1/83.4	88.8	66.5	89.9	90.5	80.3
RKD (Park et al., 2019)	66M	53.4	91.7	-	86.1	68.6	89.5	90.9	80.03
CKD (Park et al., 2021b)	66M	55.1	93.0	83.6/84.1	89.6	67.3	90.5	91.2	81.11
SFTN (Park et al., 2021a)	66M	53.6	91.5	-	85.3	68.5	89.5	90.4	79.80
TAKD (Mirzadeh et al., 2020)	66M	53.8	91.4	-	85.0	68.5	89.6	90.7	79.83
ProKT (Shi et al., 2020)	66M	54.3	91.3	-	86.3	68.4	89.7	90.9	80.15
MGSKD (Liu et al., 2022)	66M	49.1	91.7	83.3/83.9	89.8	67.9	90.3	91.2	80.00
MetaDistill (Zhou et al., 2022)	66M	58.6	92.3	-	86.8	69.4	90.4	91.0	81.42
AD-KD (Wu et al., 2023)	66M	58.3	91.9	83.4/84.2	91.2	70.9	91.2	91.2	82.45
Task-specific Logits-based Distillation									
Vanilla KD (Hinton et al., 2015)	66M	53.6	91.1	82.7/83.1	89.4	66.8	90.1	90.5	80.25
RCO (Jin et al., 2019)	66M	53.6	91.4	-	85.1	67.6	89.7	90.6	79.67
DML (Zhang et al., 2018)	66M	53.7	91.5	-	85.1	68.4	89.6	90.3	79.77
PD (Turc et al., 2019)	66M	-	91.1	82.5/83.4	89.4	66.7	89.4	90.7	-
ReAugKD (Zhang et al., 2023)	66M	59.4	92.5	-	86.3	70.4	90.7	91.2	81.75
SinKD (ours)	66M	60.2	93.1	83.8/84.2	91.3	71.1	90.5	91.3	82.92

Takeaways & Preview

- OT provides a flexible and geometry-aware loss function for comparing distributions
- Entropic regularization makes OT **differentiable, smooth, and tractable** for modern ML workflows
- Gradient formulas exist for both discrete (Eulerian/Lagrangian) and continuous measures
- Automatic differentiation enables scalable training by **integrating OT loss directly into deep learning pipelines**
- OT-based losses are widely used in generative modeling, structured prediction, and domain adaptation
- **Next time:** guest lecture by Kimia Nadjahi on Sliced OT